

# 103590450 四資四 馬茂源

1. It is known that the HIV test has only 0.1% of false positive and false negative, respectively. However, for a specific group of people, the prevalence of HIV positive rate is 0.01 %. If a person belongs to such a group and is found to be positive in the HIV test, find the probability that the person is really infected.

	true	false
test positive	0.01% * 99.9%	99.99% * 0.1%
test negative	0.01% * 0.1%	99.99% * 99.9%
	0.01%	99.99%

P(really infected | test positive) = precision =

$$\frac{TP}{TP + FP}$$

$$= 0.01\% * 99.9\% / (0.01\% * 99.9\% + 99.99\% * 0.1\%) = 0.09083469721767594 = \mathbf{9.08\%}$$

```
[1] from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.datasets import load_iris
import numpy as np
import matplotlib.pyplot as plt
import operator
```

```
[2] iris = load_iris()
print(iris.DESCR)
```

Iris Plants Database  
=====

Notes  
-----

Data Set Characteristics:

```

:Number of Instances: 150 (50 in each of three classes)
:Number of Attributes: 4 numeric, predictive attributes and the class
:Attribute Information:
  - sepal length in cm
  - sepal width in cm
  - petal length in cm
  - petal width in cm
  - class:
    - Iris-Setosa
    - Iris-Versicolour
    - Iris-Virginica
:Summary Statistics:

=====  =====  =====  =====  =====
              Min    Max    Mean     SD    Class Correlation
=====  =====  =====  =====  =====
sepal length:  4.3    7.9    5.84    0.83     0.7826
sepal width:   2.0    4.4    3.05    0.43    -0.4194
petal length:  1.0    6.9    3.76    1.76     0.9490 (high!)
petal width:   0.1    2.5    1.20    0.76     0.9565 (high!)
=====  =====  =====  =====  =====

:Missing Attribute Values: None
:Class Distribution: 33.3% for each of 3 classes.

```

**2 . UC Irvine has a large repository for various kinds of data. In this problem, you are asked to use the iris dataset (<https://archive.ics.uci.edu/ml/datasets/Iris>) to perform the experiments. Implement the k-NN classifier for the classification task. To begin one experiment, randomly draw 70 % of the instances for training and the rest for testing. Repeat the drawing and the k-NN classification 10 times and compute the average accuracy. Then, plot the curve of k versus accuracy for k = 1, 3, ..., 15. For simplicity, use the Euclidean distance in your computation.**

```
[3] N = 10 #200次的話曲線比較穩定 不過題目說10次就用10次
```

```
[4] class MyKNeighborsClassifier:

    def __init__(self, n_neighbors=3, **kwargs):
        self._k = n_neighbors
        self._X = self._y = None
        self.set_params(**kwargs)

    def get_params(self, deep=True):
        # suppose this estimator has parameters "alpha" and "recursive"
        return self.__dict__
```

```

def set_params(self, **parameters):
    for parameter, value in parameters.items():
        setattr(self, parameter, value)
    return self

def fit(self, X, y):
    self._X = X
    self._y = y

def _predict(self, x):
    distances = np.apply_along_axis(lambda x1: np.linalg.norm(x-x1),
                                    1, self._X)
    X_candidates = np.argsort(distances)[:self._k]
    y_candidates = self._y[X_candidates]
    return np.argmax(np.bincount(y_candidates))

def score(self, X, y_true):
    return accuracy_score(y_true, self.predict(X))

def predict(self, X):
    return np.apply_along_axis(lambda x: self._predict(x), 1, X)

```

```

[5] K = list(range(1, 15+1, 2))
    Knn_acc = []
    for k in K:
        # Repeat the drawing and the k-NN classification 10 times.
        acc = []

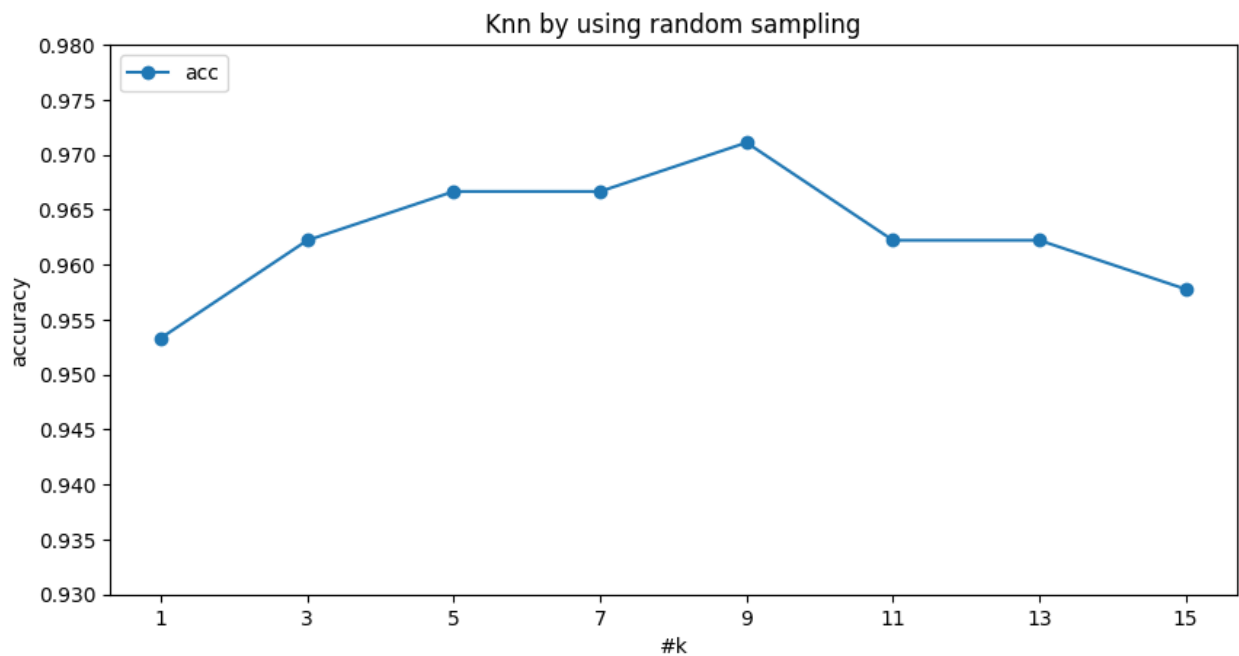
        for i in range(N):
            model = MyKNeighborsClassifier(n_neighbors=k)
            X_train, X_test, y_train, y_test = train_test_split(iris.data,
                                                                iris.target,
                                                                test_size=0.3)

            model.fit(X_train, y_train)
            acc.append(accuracy_score(y_test, model.predict(X_test)))

        Knn_acc.append(np.mean(acc))

fig = plt.figure(figsize=(10, 5), dpi=100, facecolor='white')
plt.plot(K, Knn_acc, 'o-')
plt.xticks(np.arange(1, 15+1, 2))
plt.yticks(np.arange(0.93, 0.98+0.005, 0.005))
plt.title('Knn by using random sampling')
plt.ylabel('accuracy')
plt.xlabel('#k')
plt.legend(['acc'], loc='upper left')
plt.show()

```



**3 . Following problem 2, if you do not have the test dataset (i.e., you have only the 70 % of dataset), how do you determine the optimal value of k? Use your own approach to find such a value and compare the results you have in problem 2 . Comment on your results.**

**ANS:**

**I use K-fold cross-validation**

**Note:**

如果以10次來平均的話 根本無法選擇k，Problem2跟3的曲線無法穩定，必須增加到200次以上才會穩定

不過就在10次平均下多跑幾次來看，還是可以發現大概在k>9時精準度會開始下降

```
[9] K = list(range(1, 15+1, 2))
Knn_acc_vc = []
for k in K:
    acc = []
    for i in range(N):
        model = MyKNeighborsClassifier(n_neighbors=k)
        # Assuming that I don't have test data
        X_train, _, y_train, _ = train_test_split(iris.data,
                                                    iris.target,
                                                    test_size=0.3)
        acc.append(np.mean(cross_val_score(model, X_train,
```


```

y_train, cv=5)))

Knn_acc_vc.append(np.mean(acc))

fig = plt.figure(figsize=(10, 5), dpi=100, facecolor='white')
plt.plot(K, Knn_acc_vc, 'o-', label='problem 3')
plt.plot(K, Knn_acc, '--', label='problem 2')
plt.xticks(np.arange(1, 15+1, 2))
plt.yticks(np.arange(0.93, 0.98+0.005, 0.005))
plt.title('Knn by using K-fold cross-validation')
plt.ylabel('accuracy')
plt.xlabel('#k')
plt.legend(loc='upper left')
plt.show()

```



**4 . In the class, we covered the naive Bayes classifier, but only with discrete-type features. Consult any paper to learn how to extend this approach to continuous-type features. Explain your finding as an algorithm.**

Bayesian probability

$$p(C_k | \mathbf{x}) = \frac{p(C_k) p(\mathbf{x} | C_k)}{p(\mathbf{x})}$$

When dealing with continuous data, a typical assumption is that the continuous values associated with each class are distributed according to a Gaussian distribution.

$$P(x = v|c) = \frac{1}{\sqrt{2\pi\sigma_c^2}} e^{-\frac{(v-\mu_c)^2}{2\sigma_c^2}}$$

The joint model can be expressed as

$$\begin{aligned} p(C_k | x_1, \dots, x_n) &\propto p(C_k, x_1, \dots, x_n) \\ &\propto p(C_k) p(x_1 | C_k) p(x_2 | C_k) p(x_3 | C_k) \dots \\ &= p(C_k) \prod_{i=1}^n p(x_i | C_k). \end{aligned}$$

We define classifier like this

$$\hat{y} = \underset{k \in \{1, \dots, K\}}{\operatorname{argmax}} p(C_k) \prod_{i=1}^n p(x_i | C_k).$$

**5 . Repeat problem 2 with your algorithm in problem 4. Compare the accuracy of naive Bayes classifier with the k-NN.**

```
[7] class MyGaussianNB:

    def __init__(self, **kwargs):
        self._class = {}
        self.set_params(**kwargs)

    def get_params(self, deep=True):
        return self.__dict__

    def set_params(self, **parameters):
        for parameter, value in parameters.items():
            setattr(self, parameter, value)
        return self

    def fit(self, X, y):
        for Ck in np.unique(y):
            data = X[np.where(y == Ck)]
            self._class[Ck] = {'data': data}
            self._class[Ck]['mean'] = np.mean(data, axis=0)
            self._class[Ck]['var'] = np.var(data, axis=0)
            self._class[Ck]['prior'] = (data.shape[0]
                                         / X.shape[0])

    def _likelihood(self, v, feature_index, Ck):
        var = self._class[Ck]['var'][feature_index]
        mean = self._class[Ck]['mean'][feature_index]
        exp = np.exp((-1 * (v - mean)**2) / (2 * var**2))
        return (1 / (np.sqrt(2 * np.pi * var**2))) * exp

    def _predict(self, x):
        class_prob = {Ck: None for Ck in self._class.keys()}

        for Ck, v in class_prob.items():
```

```

        prior = self._class[Ck]['prior']
        class_prob[Ck] = prior
        for i, xn in enumerate(x):
            class_prob[Ck] *= self._likelihood(xn, i, Ck)

    result = sorted(class_prob.items(),
                    key=operator.itemgetter(1),
                    reverse=True)
    return result[0][0]

def score(self, X, y_true):
    return accuracy_score(y_true, self.predict(X))

def predict(self, X):
    return np.apply_along_axis(lambda x: self._predict(x), 1, X)

```

```

[8] K = list(range(1, 15+1, 2))
    GNB_acc = []

    for i in range(N):
        GNB_model = MyGaussianNB()
        X_train, X_test, y_train, y_test = train_test_split(iris.data,
                                                            iris.target,
                                                            test_size=0.3)

        GNB_model.fit(X_train, y_train)
        GNB_acc.append(accuracy_score(y_test, GNB_model.predict(X_test)))

    print('acc:%.3f'%(np.mean(GNB_acc)))
    fig = plt.figure(figsize=(10, 5), dpi=100, facecolor='white')
    plt.plot(K, Knn_acc, 'o-', label='Knn')
    plt.plot(K, [np.mean(GNB_acc)]*len(K), label='GBN')
    plt.xticks(np.arange(1, 15+1, 2))
    plt.yticks(np.arange(0.90, 0.98+0.005, 0.005))
    plt.title('Gaussian Naive Bayes')
    plt.ylabel('accuracy')
    plt.xlabel('#k')
    plt.legend(loc='upper left')
    plt.show()

```

acc:0.944

